Lightweight and Modular Resource Leak Verification

Anonymous Author(s)

ABSTRACT

A resource leak occurs when a program allocates a resource, such as a socket or file handle, but fails to deallocate it. Resource leaks are an important, common problem that cause resource starvation, slowdowns, and crashes. Previous techniques to prevent resource leaks are unsound, imprecise, inapplicable to existing code, slow, or a combination of these.

We observe that detecting a resource leak for a variable involves three parts: 1) tracking its *must-call obligations*, 2) tracking which methods have been called on it, and 3) comparing the results of these to check if its obligations have been fulfilled. Our key insight is that these can be reduced to an accumulation problem, a class of typestate problems amenable to sound and modular checking without the need for a heavyweight, whole-program alias analysis. We developed a baseline leak checker via this approach. The precision of an accumulation analysis can be improved by computing targeted aliasing information, and we devised three novel techniques that use this capability to achieve precision in practice: a lightweight ownership transfer system; a specialized resource alias analysis; and a system to create a fresh obligation when a non-final resource field is updated.

Our approach occupies a unique slice of the design space when compared to prior approaches: it is sound and runs quickly compared to heavier-weight approaches (running in minutes on programs that a state-of-the-art approach took hours to analyze). We implemented our techniques for Java in an open-source tool called Plumber. Plumber revealed 45 real bugs in widely-deployed software. It scales well, has a manageable false positive rate (lower than the high-confidence resource leak analysis built into the Eclipse IDE), and imposes only a small annotation burden (1/2000 LoC) for developers.

1 INTRODUCTION

A resource leak occurs when some finite resource managed by the programmer is not explicitly disposed of. In an unmanaged language like C, that explicit resource might be memory; in a managed language like Java, it might be a file descriptor, a socket, or a database connection. Resource leaks continue to cause severe bugs, even in modern, heavily-used Java applications [12]. This state-of-the-practice does not differ much from two decades ago [31]. Microsoft engineers consider resource leak bugs to be one of the most significant development challenges [20]. That resource leaks remain such a serious problem despite decades of research and improvements

in languages and tooling shows that preventing them remains an urgent, difficult, open problem.

Ideally, a tool for preventing resource leaks would be:

- *applicable* to existing code with few code changes,
- *sound*, so that undetected resource leaks do not slip into the program;
- *precise*, so that developers are not bothered by excessive false positive warnings; and
- *fast*, so that it scales to real-world programs and developers can use it regularly.

Extant approaches fail at least one of these criteria. Language-based features may not apply to all uses of resource variables: Java's try-with-resources statement [21], for example, can only close resource types that implement the `java.lang.AutoCloseable` interface, and cannot handle common resource usage patterns that span multiple procedures. Heuristic bug-finding tools for leaks, built into Java IDEs including Eclipse [7] and IntelliJ IDEA [16], are fast and applicable to legacy code—but they are unsound and can be highly imprecise (section 8.3). Inter-procedural typestate or dataflow analyses [29, 35] achieve more precise results—and can find more complex bugs than bug-finders, though they usually remain unsound—but they can require hours to analyze a large-scale modern Java program, due to the costliness of precise whole-program analysis. Finally, ownership type systems [4] as employed in languages like Rust [18] can prevent nearly all resource leaks (see section 10.2), but using them would require a significant rewrite for a legacy codebase, a substantial task which is often infeasible.

The essence of finding resource leaks in a Java-like language is checking that a particular method (typically `close()`) is called on every object of a relevant class during its lifetime (i.e., before it becomes unreachable). We deem this a *must-call* property. Our key insight is that checking must-call properties is an accumulation problem. Our contribution is an *accumulation analysis* for the resource leak problem, which satisfies all four requirements: it is applicable, sound, precise, and fast.

An accumulation analysis [17] is a special-case of typestate analysis [25]. Typestate analysis attaches a finite-state machine (FSM) to each program element of a given type, and transitions the state of the FSM whenever a relevant operation—calling a method, going out of scope, etc.—is performed. In an accumulation analysis, the order of operations performed cannot change what is subsequently permitted, and executing more operations cannot add additional restrictions. If it is applicable, accumulation analysis provides a major benefit over general typestate analysis: it does not require a whole-program alias analysis for soundness. Hence, we can build a sound, modular must-call checker without doing *any* whole-program alias analysis. Accumulation also enables a clean design, as we can conceptually factor our checker into a sound core, and then enhance it with targeted alias reasoning to reduce false positives.

Unlike recent work on accumulation [17], we found that checking resource leaks precisely requires significantly more reasoning about aliasing patterns. We found several key patterns to handle:

- passing of resources via parameters or returns;
- storing of resources in final fields, in a standard RAII manner [26];
- wrapper types, which share their must-call obligations with one of their fields; and,
- storing resources in non-final fields, which might be lazily initialized or written more than once.

We devised three key techniques to handle these patterns modularly and achieve precision in practice:

- a lightweight, heuristic ownership transfer system. This system indicates which reference is responsible for resolving a must-call obligation. Unlike typical ownership type systems, our approach does not impact the privileges of non-owning references.
- resource aliasing, for cases in which a resource's must-call obligations can be resolved by closing one of two or more references.
- a system for creating new obligations at locations other than the constructor, which allows Plumber to handle lazy or multiple initialization.

Variants of (some of) these ideas exist in previous work. We bring them together in a general, modular manner, with full verification and the ability for programmers to easily extend checking to their own types and must-call properties.

Our contributions are:

- the insight that the resource leak problem is an accumulation problem, and type systems designed to take advantage of this fact (section 3).
- three innovations that improve the precision of our type systems: a lightweight ownership transfer system (section 4), a lightweight resource-alias tracking analysis (section 5), and a system for handling lazy or multiple initialization (section 6).
- an open-source implementation for Java, called Plumber (section 7).
- an extensive empirical evaluation: case studies on heavily-used Java programs that handle many resources (section 8.1), an ablation study that shows the contributions of each innovation to Plumber's precision (section 8.2), and a comparison to other state-of-the-art approaches that demonstrates the unique strengths of our approach (section 8.3).

2 BACKGROUND ON PLUGGABLE TYPES

Sections 3.1 and 3.2 describe *pluggable type systems* [10] that are layered on top of the type system of the host language. Types in a pluggable type system are composed of two parts: a *type qualifier* and a base type. The type qualifier is the part of the type that is unique to the pluggable type system; the base type is a type from the host language. Our implementation is for Java (see section 7), so we use the Java syntax for type qualifiers: "@" before a type indicates that it is a type qualifier, and a type without "@" is a base type. This paper sometimes omits the basetype when it is obvious from context.

A type system checks programmer-written types. Our system requires the programmer to write types on method signatures, but within method bodies it uses flow-sensitive type refinement,

```
Socket s = null;
try {
  s = new Socket(myHost, myPort);
} catch (Exception e) { // do nothing
} finally {
  if (s != null) {
    s.close();
  }
}
```

Figure 1: A safe use of a `Socket` resource.

$$\text{@MustCallUnknown} = \top$$
$$\uparrow$$
$$\text{@MustCall(\{"a", "b"\})}$$

@MustCall({"a"})   @MustCall({"b"})

$$\text{@MustCall(\{\})} = \bot$$

Figure 2: Part of the `MustCall` type hierarchy for representing which methods must be called; the full hierarchy is a lattice of arbitrary size. If an expression's type has qualifier @MustCall({"a", "b"}), then the methods "a" and "b" might need to be called before the expression is deallocated. Arrows represent subtyping relationships.

a dataflow analysis that performs type inference. This permits an expression to have different types on different lines of the program.

3 LEAK DETECTION VIA ACCUMULATION

This section presents a sound, modular, accumulation-based resource leak checker ("Plumber"). Sections 4–6 enhance its precision.

Plumber is composed of three cooperating analyses:

(1) a taint tracking type system (section 3.1) computes a conservative *overapproximation* of the set of methods that might need to be called on each expression in the program.
(2) an accumulation type system (section 3.2) computes a conservative *underapproximation* of the set of methods that are actually called on each expression in the program.
(3) a dataflow analysis (section 3.3) determines whether the methods that might need to be called on each expression are always invoked before the expression goes out of scope, by checking consistency of the results of the two above-mentioned type systems; it issues errors when this condition does not hold.

This section uses fig. 1 as a motivating example. It shows a safe use of a `Socket`—a resource that must be closed before it is deallocated.

3.1 A type system for must-call obligations

The Must Call type system tracks which methods might need to be called on a given object before the object is deallocated. This type system is general—it is not specific to resource leaks.

The Must Call type system supports two qualifiers: @MustCall and @MustCallUnknown. The @MustCall qualifier's arguments are the methods that the annotated type must call. The declaration @MustCall({"a"}) Object obj means that before obj is deallocated, obj.a() might need to be called. Plumber conservatively requires all these methods to be called, and it issues a warning if they are not.

For example, consider fig. 1. The expression `null` has type `@MustCall({})`—it has no obligations to call particular methods—so `s` has that type after its initialization. The `new` expression has type `@MustCall("close")`, and therefore `s` has that type after the assignment. At the start of the `finally` block, where both values for `s` flow, the type of `s` is their least upper bound, which is `@MustCall("close")`.

Part of the type hierarchy appears in fig. 2. All types are subtypes of `@MustCallUnknown`. The subtyping relationship for `@MustCall` types is:

$$\frac{A \subseteq B}{@MustCall(A) \sqsubseteq @MustCall(B)}$$

The default type qualifier is `@MustCall({})` for base types without a programmer-written type qualifier.[1] Our implementation provides JDK annotations which require that every `Closeable` object must have the `close()` method called before it is deallocated.

3.2 A type system for called methods

The Called Methods type system tracks a conservative underapproximation of which methods have been called on an object. It is an extension of a similar system from prior work [17]. The primary difference in our version is that a method is considered called even if it throws an exception—a necessity in Java because the `close()` method in `java.io.Closeable` is specified to possibly throw an `IOException`. In the prior work, a method was only considered "called" when it terminated successfully. The remainder of this section is a brief summary of the prior work [17].

The checker is an accumulation analysis whose accumulation qualifier is `@CalledMethods`. The type `@CalledMethods(A) Object` represents an object on which the methods in the set A have definitely been called; other methods not in A might also have been called. The subtyping rule is:

$$\frac{B \subseteq A}{@CalledMethods(A) \sqsubseteq @CalledMethods(B)}$$

The top type is `@CalledMethods({})`. The qualifier `@CalledMethodsBottom` is a subtype of every `@CalledMethods` qualifier.

Thanks to flow-sensitive type refinement, Called Methods types are inferred within method bodies. In fig. 1 the type of `s` is initially `@CalledMethods({})`, but it transitions to `@CalledMethods("close")` after the call to `close`.

3.3 Consistency checking

Given `@MustCall` and `@CalledMethods` types, the Must Call Consistency Checker ensures that the `@MustCall` methods for each object are always invoked before it becomes unreachable, via an intraprocedural dataflow analysis. Here, we present a simple, sound version of the analysis, with limited reasoning about aliasing. Sections 4–6 describe enhancements to this basic approach.

Language. For simplicity, we present the analysis over a simple assignment language in three-address form. An expression e in the language is `null`, a variable `p`, a field read `p.f`, or a method call `m(p1,p2,...)` (constructor calls are treated as method calls). A statement s takes one of three forms: `p = e`, where e is an expression; `p.f = p'`, for a field write; or `return p`. Methods are represented by a control-flow graph (CFG) where nodes are statements

[1]For unannotated local variable types, flow-sensitive type refinement infers a qualifier.

Algorithm 1 Finding unfulfilled `@MustCall` obligations in a method. Algorithm 2 defines helper functions.

1: **procedure** FindMissedCalls(*CFG*)
2:   // *D maps each statement s to a set of dataflow facts reaching*
3:   // *s. Each fact is of the form* $\langle P, e \rangle$*, where P is a set of variables*
4:   // *that must-alias e and e is an expression with a nonempty*
5:   // *must-call obligation.*
6:   $D \leftarrow$ InitialObligations(*CFG*)
7:   **while** D has not reached fixed point **do**
8:     **for** $s \in CFG.statements$, $\langle P, e \rangle \in D(s)$ **do**
9:       **if** s is *exit* **then**
10:         report a must-call violation for e
11:       **else if** ¬MCSatisfiedAfter(P, s) **then**
12:         $kill \leftarrow s$ assigns a variable ? $\{s.LHS\} : \emptyset$
13:         $gen \leftarrow$ CreatesAlias(P, s) ? $\{s.LHS\} : \emptyset$
14:         $N \leftarrow (P - kill) \cup gen$
15:         $\forall t \in CFG.succ(s) . D(t) \leftarrow D(t) \cup \langle N, e \rangle$
16: **procedure** InitialObligations(*CFG*)
17:   $D \leftarrow \{s \mapsto \emptyset \mid s \in CFG.statements\}$
18:   **for** $p \in CFG.formals$, $t \in CFG.succ(CFG.entry)$ **do**
19:     **if** HasObligation(p) **then**
20:       $D(t) \leftarrow D(t) \cup \langle\{p\}, p\rangle$
21:   **for** $s \in CFG.statements$ of the form `p = m(p1, p2, ...)` **do**
22:     $\forall t \in CFG.succ(s) . D(t) \leftarrow D(t) \cup$ FactsFromCall(s)
23:   **return** D

Algorithm 2 Helper functions for algorithm 1. Except for MCAfter and CMAfter, all functions will be replaced with more sophisticated versions in sections 4–6.

1: // *Does e introduce a must-call obligation to check?*
2: **procedure** HasObligation(e)
3:   **return** e has a declared `@MustCall` type
4: // *s must be a call statement p = m(p1, p2, ...)*
5: **procedure** FactsFromCall(s)
6:   $p \leftarrow s.LHS$, $c \leftarrow s.RHS$
7:   **return** HasObligation(c) ? $\{\langle\{p\}, c\rangle\} : \emptyset$
8: // *Is the must-call obligation for P satisfied after s?*
9: **procedure** MCSatisfiedAfter(P, s)
10:   **return** $\exists p \in P$. MCAfter(p, s) $\subseteq$ CMAfter(p, s)
11: // *Does s introduce a must alias for a var in P?*
12: **procedure** CreatesAlias(P, s)
13:   **return** $\exists q \in P$. s is of the form `p = q`
14: **procedure** MCAfter(p, s)
15:   **return** methods in `@MustCall` type of p after s
16: **procedure** CMAfter(p, s)
17:   **return** methods in `@CalledMethods` type of p after s

---

and edges indicate possible control flow. We elide control-flow predicates as the consistency checker is path-insensitive.

For a method CFG, *CFG.statements* is the statements, *CFG.formals* is the formal parameters, *CFG.entry* is its entry node, *CFG.exit* is its exit node, and *CFG.succ* is its successor relation. For a statement s of the form `p = e`, $s.LHS = p$ and $s.RHS = e$.

```
s = new Socket(...); // 1
if (...) {
    s = null; // 2
} else {
    t = s; // 3
    close(t); // 4
}
```
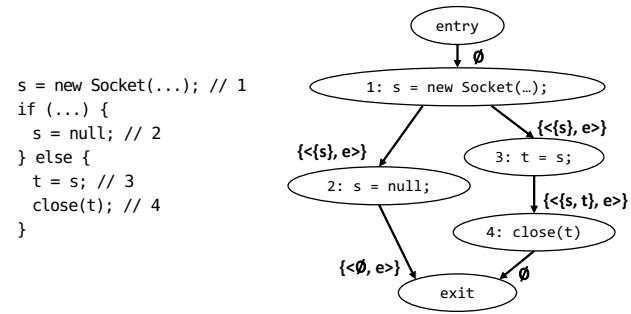

Figure 3: Example code and CFG for illustrating algorithm 1. "e" is "new Socket(...)".

Pseudocode. Algorithm 1 gives pseudocode for the basic version of our checker, with helper functions in algorithm 2. At a high level, the dataflow analysis computes a map D from each statement s in a CFG to a *set* of facts of the form $\langle P, e \rangle$, where P is a set of variables and e is an expression. The meaning of D is as follows: if $\langle P, e \rangle \in D(s)$, then e has a declared @MustCall type, and all variables in P are *must aliases* for the value of e at the program point before s. Computing a set of must aliases is useful since any must alias may be used to satisfy the must-call obligation of e. Using D, the analysis finds any e that does not have its @MustCall obligation fulfilled, and reports an error.

Algorithm 1 proceeds as follows. Line 6 invokes INITIALOBLIGA-TIONS to initialize D. Only formal parameters or method calls can introduce obligations to be checked (not reads of local variables or fields). The fixed-point loop iterates over all facts $\langle P, e \rangle$ present in any $D(s)$ (our implementation uses a worklist for efficiency). If s is the exit node (line 9), the obligation for e has not been satisfied, and an error is reported. Otherwise, the algorithm checks if the obligation for e is satisfied after s (line 11). For the basic checker, MCSATISFIEDAFTER in algorithm 2 checks whether there is some $p \in P$ such that after s, the set of methods in p's @MustCall type are contained in the set of methods in its @CalledMethods type; if true, all @MustCall methods have already been invoked. This check uses the inferred flow-sensitive @MustCall and @CalledMethods qualifiers described in sections 3.1 and 3.2.

If the obligation for e is not yet satisfied, the algorithm propagates the fact to successors with an updated set N of must aliases. N is computed in a standard gen-kill style on lines 12–14. The kill set simply consists of whatever variable (if any) appears on the left-hand side of s. The gen set is computed by checking if s creates a new must alias for some variable in P, using the CREATESALIAS routine. Since our analysis is accumulation, CREATESALIAS could simply return false without impacting soundness. In algorithm 2, CREATESALIAS handles the case of a variable copy where the right-hand side is in P. (Section 5 presents more sophisticated handling.) Finally, line 15 propagates the new fact to successors. The process continues until D reaches a fixed point.

Example. To illustrate our analysis, fig. 3 shows a simple program (irrelevant details elided) and its corresponding CFG. The CFG shows the dataflow facts propagated along each edge. For initialization, statement 1 introduces the fact $\langle \{s\}, e \rangle$ (where e is the new Socket(...) call) to $D(2)$ and $D(3)$. At statement 2, s is killed,

causing $\langle \emptyset, e \rangle$ to be added to $D(exit)$. This leads to an error being reported for statement 1, as the socket is not closed on this path. Statement 3 creates a must alias t for s, causing $\langle \{s, t\}, e \rangle$ to be added to $D(4)$. For statement 4, MCSATISFIEDAFTER($\{s, t\}$, close(t)) holds, so no facts are propagated from 4 to *exit*.

At this point, we have described a sound checker for @MustCall obligations. Subsequent sections make the checker more precise.

4 LIGHTWEIGHT OWNERSHIP TRANSFER

Section 3 describes a sound accumulation-based checker for resource leaks. However, that checker often encounters false positives in cases where an @MustCall obligation is satisfied in another procedure, via parameter passing, return values, or object fields. Consider the following code that safely closes a Socket:

```
void example(String myHost, int myPort) {
    Socket s = new Socket(myHost, myPort);
    closeSocket(s);
}
void closeSocket(@Owning @MustCall("close") Socket t) {
    t.close();
}
```

The closeSocket() routine takes ownership of the socket—that is, it takes responsibility for closing it. The checker described by section 3 would issue a false positive on this code, because it would warn when s goes out of scope in example().

This section describes a *lightweight ownership transfer* technique for reducing false positives in such cases. Programmers write annotations such as @Owning that transfer an obligation: they take away an obligation in one place in the program and create an obligation somewhere else. Programmer annotations *cannot* introduce any checker unsoundness; at worst, incorrect @Owning annotations will cause false positive warnings. Unlike an ownership type system like Rust's (see section 10.2), lightweight ownership transfer imposes no restrictions on what operations can be performed through an alias, and hence has a minimal impact on the programming model.

4.1 Ownership transfer

@Owning is a declaration annotation, not a type qualifier; it can be written on a declaration such as a parameter, return, field, etc., but not on a type. A pseudo-assignment to an @Owning lvalue transfers the right-hand side's @MustCall obligation. More concretely, in the Must Call Consistency Checker (Section 3.3), at a pseudo-assignment to an lvalue with an @Owning annotation, the right-hand side's @MustCall obligation is treated as satisfied.

The MCSATISFIEDAFTER(P, s) and HASOBLIGATION(e) procedures of algorithm 2 are enhanced for ownership transfer as follows:

procedure MCSATISFIEDAFTER(P, s)
    return $\exists p \in P.$ MCAFTER(p, s) $\subseteq$ CMAFTER(p, s)
             $\vee$ (s is return p $\wedge$ OWNINGRETURN(CFG))
             $\vee$ PASSEDASOWNINGPARAM(s, p)
             $\vee$ (s is q.f = p $\wedge$ f is @Owning)

procedure HASOBLIGATION(e)
    return e has a declared @MustCall type and e's declaration is @Owning

procedure OWNINGRETURN(CFG)
    return CFG's return declaration is @Owning

procedure PassedAsOwningParam(s,p)
    return s passes p to an @Owning parameter of its callee

Section 4.2 discusses checking of @Owning fields.

Constructor returns are always @Owning. Plumber's default for unannotated method returns is @Owning, and for unannotated parameters and fields is @NotOwning. These assumptions coincide well with coding patterns we observed in practice, reducing the annotation burden for programmers. Further, this treatment of parameter and return types ensures sound handling of unannotated third-party libraries: any object returned from such a library will be tracked by default, and the checker never assumes that passing an object to an unannotated library will satisfy its obligations.

4.2 Final owning fields

Additional class-level checking is required for @Owning fields, as the code satisfying their @MustCall obligations nearly always spans multiple procedures. We first discuss the case of final fields,[2] which cannot be overwritten after initialization of the enclosing object. Here, our checking enforces the "resource acquisition is initialization (RAII)" programming idiom [26]; a distinguished "destructor"-like method d() must ensure the field's @MustCall obligation is satisfied, and the enclosing class must have a @MustCall("d") obligation to ensure the destructor is called.

More formally, consider a final @Owning field f declared in class C, where f has type @MustCall("m"). To verify that f's @MustCall obligation is satisfied, Plumber checks the following conditions:

(1) All C objects must have a type @MustCall("d") for some method C.d().
(2) C.d() must always invoke this.f.m(), thereby satisfying f's @MustCall obligation.

In this manner, Plumber modularly verifies the @MustCall obligation of f. Condition 2 is checked by requiring an appropriate @EnsuresCalledMethods postcondition annotation on C.d(), which is then enforced by the Called Methods Checker [17].

While this scheme suffices for verifying final fields, we observed cases in our benchmarks of non-final fields with @MustCall obligations. Checking non-final fields requires additional techniques, as the checker must ensure that overwriting the field is safe. Section 6.1 describes our novel handling of non-final fields.

5 RESOURCE ALIASING

This section introduces a sound, lightweight, specialized must-alias analysis that tracks *resource alias* sets—sets of pointers that definitely correspond to the same underlying system resource. Closing one alias also closes the others. Thus, Plumber can avoid issuing false positive warnings about resources that have already been closed through a resource alias. Prior work on accumulation analyses showed that limited, specialized forms of alias analysis can provide precision to an accumulation analysis where needed, without incurring the cost of a whole-program analysis [17]; our resource alias analysis is an example.

---

[2]Plumber treats all static fields as non-owning, meaning that no assignment to one can satisfy a must-call obligation. In our case studies, we did not observe any assignments of expressions with non-empty must-call obligations to static fields. We leave handling owning static fields to future work.

5.1 Wrapper types

Java programs extensively use *wrapper types*. For example, the Java BufferedOutputStream wrapper adds buffering to some other delegate OutputStream, which may or may not represent a resource that needs closing. The wrapper's close() method invokes close() on the delegate. Wrapper types introduce two additional complexities for @MustCall checking:

(1) If a wrapped object has no @MustCall obligation, the corresponding wrapper object should also have no obligation.
(2) Satisfying the obligation of *either* the wrapped object or the wrapper object is sufficient.

For example, if a BufferedOutputStream b wraps a stream with no underlying resource (e.g., a ByteArrayOutputStream), b's @MustCall obligation should be empty, as b has no resource of its own. By contrast, if b wraps a stream managing a resource, like a FileOutputStream f, then close() must be invoked on *either* b or f.

Previous work has shown that reasoning about wrapper types is required to avoid excessive false positive and duplicate reports [7, 29]. Wrapper types in earlier work were handled with hard-coded specifications of which library types are wrappers, and heuristic clustering to avoid duplicate reports for wrappers [29].

Our technique handles wrapper types more generally by tracking *resource aliases*. Two references r_1 and r_2 are resource aliases if r_1 and r_2 are must-aliased pointers or if satisfying r_1's @MustCall obligation also satisfies r_2's obligation, and vice-versa.

Introducing resource aliases. To indicate where an API method creates a resource-alias relationship between distinct objects, the programmer writes a pair of @MustCallAlias qualifiers: one on a parameter of a method, and another on its return type. For example, one constructor of BufferedOutputStream is:

```
@MustCallAlias BufferedOutputStream(@MustCallAlias OutputStream arg0);
```

@MustCallAlias annotations are verified, not trusted; see section 5.3. At call sites to @MustCallAlias methods, there are two effects.

First, the must-call type of the method call's return value is the same as that of the @MustCallAlias argument. If the type of the argument has no must-call obligations (like a ByteArrayOutputStream), the returned wrapper will not have must-call obligations, either.

Second, the Must Call Consistency Checker (Section 3.3) treats the @MustCallAlias parameter and return as aliases. For our section 3.3 pseudocode, this version of CreatesAlias from algorithm 2 handles resource aliases:

procedure CreatesAlias(P, s)
    return $\exists q \in P$. s is of the form p = q
                $\lor$ IsMustCallAliasParam(s, q)
procedure IsMustCallAliasParam(s, p)
    return s passes p to a @MustCallAlias parameter of its callee

5.2 Beyond wrapper types

@MustCallAlias can also be employed in scenarios beyond direct wrapper types, a capability not present in previous work on resource leak detection. In certain cases, a resource gets shared between objects via an intermediate object that cannot directly close the resource. For example, java.io.RandomAccessFile (which must be closed) has a method getFd() that returns a FileDescriptor

object for the file. This file descriptor cannot be closed directly—it has no `close()` method. However, the descriptor can be passed to a wrapper stream such as `FileOutputStream`, which if closed satisfies the original must-call obligation. By adding `@MustCall-Alias` annotations to the `getFd()` method, our technique can verify code like the below (adapted from Apache Hadoop [27]):

```
RandomAccessFile file = new RandomAccessFile(myFile, "rws");
FileInputStream in = null;
try {
  in = new FileInputStream(file.getFD());
  // do something with in
  in.close();
} catch (IOException e){
  file.close();
}
```

Because the must-call obligation checker (section 3.1) treats `@Must-CallAlias` annotations polymorphically, regardless of the associated base type, Plumber can verify that the same resource is held by the `RandomAccessFile` and the `FileInputStream`, even though it is passed via a class without a `close()` method. This capability enabled verifying multiple code patterns in our case studies.

5.3 Verification of `@MustCallAlias`

A pair of `@MustCallAlias` annotations on a method or constructor `m`'s return type and its parameter `p` can be verified if either of the following holds:

(1) `p` is passed to another method or constructor in a `@MustCall-Alias` position, and `m` returns that method's result, or the call is a `super()` constructor call annotated with `@MustCallAlias`.
(2) `p` is stored in an `@Owning` field of the enclosing class. (`@Owning` field verification is described in sections 4.2 and 6.1.)

These verification procedures permit a programmer to soundly specify a resource-aliasing relationship in their own code, a capability that was not present in prior work that relied on a hard-coded list of wrapper types; we used this multiple times in our case studies.

6 CREATING NEW OBLIGATIONS

Every constructor of a class that has must-call obligations implicitly creates obligations for the newly-created object. However, non-constructor methods may also create obligations when re-assigning non-final, owning fields or allocating new system-level resources. A post-condition annotation, `@CreateObligation`, indicates for which expression an obligation is created.

At each call to a method annotated as `@CreateObligation(expr)`, Plumber removes any inferred Called Methods information about `expr`, reverting to the declared type of `expr`. Since a `@CreateObli-gation` annotation can only increase its target's obligations, no verification for them is needed: placing them anywhere in the program is sound (but may reduce precision).

When checking a call to a method annotated as `@CreateObli-gation(expr)`, the Must Call Consistency Checker (1) treats the `@MustCall` obligation of `expr` as *satisfied*, and (2) creates a fresh obligation to check. In terms of our section 3.3 pseudocode, we can update the FactsFromCall and MCSatisfiedAfter procedures of algorithm 2 to handle `@CreateObligation` as follows ([. . .] stands for the cases shown previously, including those in section 4.1):

procedure FactsFromCall(s)

$p \leftarrow s.LHS, c \leftarrow s.RHS$
return $\{\langle\{p_i\}, c\rangle \mid p_i \in \text{COTargets}(c)\}$
       $\cup (\text{HasObligation}(c) ? \{\langle\{p\}, c\rangle\} : \emptyset)$
procedure MCSatisfiedAfter(P, s)
    return $\exists p \in P. [\ldots] \lor p \in \text{COTargets}(s)$
procedure COTargets(c)
    return $\{p_i \mid p_i$ passed to a `@CreateObligation` target for c's callee $\}$

This change is sound: the checker creates a new obligation for calls to `@CreateObligation` methods, and the must-call obligation checker (section 3.1) ensures the `@MustCall` type for the target will have a *superset* of any methods present before the call. There is an exception to this check: if an `@CreateObligation` method is invoked within a method that has an `@CreateObligation` annotation with the same target—imposing the obligation on its caller—then the new obligation can be treated as satisfied immediately.

6.1 Non-final, owning fields

`@CreateObligation` allows Plumber to verify uses of non-final fields that contain a resource, even if they are re-assigned. Consider the following example:

```
@MustCall("close") // default qualifier for uses of SocketContainer
class SocketContainer {
  private @Owning Socket sock;
  public SocketContainer() { sock = ...; }
  void close() { sock.close() };
  @CreateObligation("this")
  void reconnect() {
    if (!sock.isClosed()) {
      sock.close();
    }
    sock = ...;
  }
}
```

In the lifetime of a `SocketContainer` object, `sock` might be re-assigned arbitrarily many times: once at each call to `reconnect()`. This code is safe, however: `reconnect()` ensures that `sock` is closed before re-assigning it.

Plumber must enforce two new rules to ensure that re-assignments to non-final, owning fields like `sock` in the example above are sound:

- any method that re-assigns a non-final, owning field of an object must be annotated with an `@CreateObligation` annotation that targets that object.
- when a non-final, owning field f is re-assigned at statement s, its inferred `@MustCall` obligation must be contained in its `@CalledMethods` type at the program point before s.

The first rule ensures that `close()` is called after the last call to `reconnect()`, and the second rule ensures that `reconnect()` safely closes `sock` before re-assigning it. Because calling an `@CreateObli-gation` method like `reconnect()` resets called-methods inference, calls to `close` before the last call to `reconnect()` are disregarded.

6.2 Unconnected sockets

`@CreateObligation` can also handle cases where object creation does not allocate a resource, but the object will allocate a resource later in its lifecycle. Consider the no-argument constructor to `java.net.Socket`. This constructor does not allocate an operating

Table 1: Verifying the absence of resource leaks in case studies. Throughout, "LoC" is lines of non-comment, non-blank Java code. "Resources" is the number of resources created by the program. "Annos." is number of manually-written annotations to specify existing methods. "Code changes" is the number of distinct changes to program text we made, not including changes that will be erased at compile time (such as annotations or warning suppressions). "TPs" is true positives. "FPs" is false positives, where the our analysis could not guarantee that the call was safe, but manual analysis revealed that no run-time failure was possible.

Project:module	LoC	Resources	Annos.	Code changes	TPs	FPs	Wall-clock time
apache/zookeeper:zookeeper-server	45,248	232	98	5	12	47	1m 24s
apache/hadoop:hadoop-hdfs-project/hadoop-hdfs	151,595	366	129	7	22	37	16m 21s
apache/hbase:hbase-server, hbase-client	220,828	62	35	1	11	27	7m 45s
Total	417,671	660	262	13	45	111	-

system-level socket, but instead just creates the container object, which permits the programmer to e.g. set options which will be used when creating the physical socket. When such a `Socket` is created, it initially has no must-call obligation; it is only when the `Socket` is actually connected via a call to a method such as `bind()` or `connect()` that the must-call obligation is created.

If all `Socket`s are treated as `@MustCall({"close"})`, a false positive would be reported in code such as the below, which operates on an unconnected socket (simplified from real code in Apache Zookeeper [28]):

```
static Socket createSocket() {
  Socket sock = new Socket();
  sock.setSoTimeout(...);
  return sock;
}
```

The call to `setSoTimeout` can throw a `SocketException` if the socket is actually connected when it is called. Using `@CreateObligation`, however, Plumber can soundly show that this socket is not connected: the type of the result of the no-argument constructor can be treated as `@MustCall({})`, and `@CreateObligation` annotations on the methods that actually allocate the socket—`connect()` or `bind()`—enforce that as soon as the socket is actually open, it is treated as `@MustCall("close")`.

7 IMPLEMENTATION

We implemented Plumber on top of the Checker Framework [22], an industrial-strength framework for building pluggable type systems for Java. The checkers which propagate and infer `@MustCall` and `@CalledMethods` annotations are implemented directly as Checker Framework checkers. The Must Call Consistency Checker (algorithm 1) is implemented as a post-analysis pass over the control-flow graph produced by the Checker Framework's dataflow analysis, which is invoked when the other two checkers terminate. The framework provides the checkers with flow-sensitive local type inference, support for Java generics and qualifier polymorphism, and other conveniences. Our implementation is open-source and available at https://github.com/fse-main-307/plumber.

8 EVALUATION

Our evaluation has three parts:

- case studies on open-source projects, which show that our approach is scalable and finds real bugs (section 8.1).
- an evaluation of the contributions of lightweight ownership, resource aliasing, and obligation creation (section 8.2).
- a comparison to previous leak detectors: both heuristic bug-finding and heavy-weight whole-program analysis (section 8.3).

Table 2: The total number of annotations that we wrote.

Annotation	Count
`@Owning` and `@NotOwning`	113
`@EnsuresCalledMethods`	43
`@MustCall`	41
`@MustCallAlias`	36
`@CreateObligation`	29

Results. All code and data for the experiments described in this section, including Plumber's implementation, our experimental machinery, and the annotated versions of our case study programs, are publicly available at https://github.com/fse-main-307.

8.1 Case studies on open-source projects

We selected 3 popular open-source projects that were analyzed by prior work [35]. For each, we selected and analyzed one or two modules with many uses of leakable resources. We used the latest version of the source code that was available when we began.

Our methodology was: (1) We modified the build system to run our analysis on the module, analyzing uses of resource classes that are defined in the JDK. It also reports the maximum possible number of resources that could be leaked: each obligation at a formal parameter or method call. (2) We manually annotated each program with must-call, called-methods, and ownership annotations (see section 8.1.2). (3) We iteratively ran the analysis to correct our annotations. We measured the run time as the median of 5 trials on a machine with an Intel Core i7-10700 CPU running at 2.90GHz and 64GiB of RAM. Our analysis is embarrassingly parallel, but our implementation is single-threaded because javac is single-threaded. (4) We manually categorized each warning as revealing a real resource leak (a true positive or TP) or as correct code that our system is unable to prove correct (a false positive or FP).

Table 1 summarizes the results. Plumber found multiple serious resource leak bugs in every program. Plumber issues more false positives than true positives in each program, but the number is small enough to be examined by a single developer in a few hours. This is a small price to pay for knowing that the program is free of resource leaks. The annotations in the program are also a benefit: as a form of machine-checked documentation, they express the programmer's intent and, unlike traditional comments, cannot become out-of-date if the checker is passing.

8.1.1 True and false positive examples. This section describes some examples of warnings reported by Plumber in our case studies.

Figure 4 contains code from Hadoop. If an IO error occurs any time between the allocation of the `FileInputStream` in the first line of the method and the `return` statement at the end—for example, if

```
813   public InputStream getInputStreamForSection(
814       FileSummary.Section section, String compressionCodec)
815       throws IOException {
816     FileInputStream fin = new FileInputStream(filename);
817     FileChannel channel = fin.getChannel();
818     channel.position(section.getOffset());
819     InputStream in = new BufferedInputStream(new LimitInputStream(fin,
820         section.getLength()));
821     in = FSImageUtil.wrapInputStreamForCompression(conf,
822         compressionCodec, in);
823     return in;
824   }
```

Figure 4: A true positive that Plumber found in Hadoop. Our pull request to fix this bug was accepted by Hadoop's developers.

```
Optional<ServerSocket> createServerSocket(...) {
  ServerSocket serverSocket;
  try {
    if (...) {
      serverSocket = new ServerSocket();
      serverSocket.setReuseAddress(true);
      serverSocket.bind(...);
      return Optional.of(serverSocket);
    }
  } catch (IOException e) {
    // log an error
  }
  return Optional.empty();
}
```

Figure 5: Code from the ZooKeeper case study that causes Plumber to issue a false positive.

`channel.position(section.getOffset())` throws an `IOException`, as it is specified to do—then the only reference to the stream is lost. Hadoop's developers assigned this bug a priority of "Major" and accepted our patch.[3] One developer suggested using a try-with-resources statement instead of our patch (which catches the exception and closes the stream), but we pointed out that the file needs to remain open if no error occurs so that it can be returned.

The most common false positive pattern (with 15 instances) was caused by a generic container object like `java.util.Optional` taking ownership of a resource, such as the example in fig. 5. Our lightweight ownership system does not support transferring ownership to generic parameters, so Plumber issues an error when `Optional.of` is returned. In this case, the use of the `Optional` class complicates the code; if `Optional` was replaced by `null`, Plumber could verify this code. We leave expanding the lightweight ownership system to support Java generics as future work.

Most false positives involve unique coding patterns. One example is a series of catch statements that each sets an `error` boolean to `true`, and a `finally` statement that closes the relevant socket if `error` was true. Plumber reports an error because it does not reason about path conditions, and handling such reasoning is future work.

8.1.2 Annotations and code changes. We wrote less than one annotation per 2,000 lines of code (table 2).

We also made 13 small, semantics-preserving changes to the programs to reduce false positives from our analysis. In 8 places, we added `final` to a field; this allows our checker to verify it without using the rules for non-final owning fields given in section 6, which are stricter. In 3 of those, we also removed assignments to the field

---

3URL removed for anonymity.

Table 3: The contribution of the lightweight ownership, resource aliasing, and obligation creation features in reducing false positives. Each entry is the number of extra false positive warnings reported by the variant with the given feature disabled on the given project.

Project	without LO	without RA	without CO
apache/zookeeper	66	97	10
apache/hadoop	95	113	0
apache/hbase	52	57	0
Total	213	267	10

after it was closed whose right-hand side was `null`; in 1 other we added an `else` clause in the constructor that assigned the field a `null` value. In 3 places, we re-ordered two statements to remove an infeasible control-flow-graph edge. In 2 places, we extracted an expression into a local variable, permitting flow-sensitive reasoning or targetting by a `@CreateObligation` annotation.

8.2   Evaluating our enhancements

The base analysis of section 3 produces significantly more false positives, because it lacks lightweight ownership (section 4), resource aliasing (section 5), and obligation creation (section 6). Each contributes to Plumber's precision by eliminating false positives. To evaluate the contribution of each enhancement, we individually disabled each feature and re-ran the experiments of section 8.1. Since all variants are sound (no false negatives), any difference in warnings is a false positive that is prevented by the feature.

Table 3 shows that each of lightweight ownership and resource aliases prevents more false positive errors than the total number of remaining false positives on each benchmarks—showing that removing either would make our technique produce an unreasonable number of false positives. The system for creating new obligations at points other than constructors reduces false positives by a smaller amount: non-final, owning field re-assignments are rare; and we encountered the unconnected socket pattern described in section 6.2 only in ZooKeeper. Nevertheless, this feature allows our tool to handle an important, if less common, coding pattern.

8.3   Comparison to other tools

Our approach represents a novel point in the design space of resource leak checkers. This section compares our approach with two other modern tools that detect resource leaks:

- The analysis built into the Eclipse Compiler for Java (ecj), which is the default approach for detecting resource leaks in the Eclipse IDE [7]. We used version 4.18.0.
- Grapple [35] represents a significant, recent improvement in the scalability of typestate-based tools that require a whole-program alias analysis.

In brief, both tools are unsound. Both tools are applicable to legacy code. Eclipse is very fast (nearly instantaneous) and suffers 67–98% false positives. According to its authors, Grapple is precise (13% false positive rate) but orders of magnitude slower than Plumber. Different users can select whichever tool matches their priorities.

8.3.1 Eclipse. The Eclipse analysis is a simple dataflow analysis augmented with heuristics. Since it is tightly integrated with the

Table 4: The Grapple tool's performance; reproduced from [35].

Project	TPs	FPs	Run time
ZooKeeper	6	0	01h 07m 02s
HDFS	5	2	01h 54m 52s
HBase	15	2	33h 51m 59s
Total	26	4	-

compiler, it scales well and runs quickly. It has heuristics for ownership, resource wrappers, and resource-free closeables, among others; these are all hard-coded into the analysis and cannot be adjusted by the user. It does not support annotations to express specifications that differ from its defaults. It supports two levels of analysis: detecting high-confidence resource leaks and detecting "potential" resource leaks (a superset of high-confidence resource leaks).

We ran Eclipse's analysis on the same versions of our case study programs that we ran Plumber on in section 8.1, and examined the same modules. It is easy to apply to legacy code and it is fast—nearly instantaneous once Eclipse has loaded the project.

In "high-confidence" mode on the three projects, Eclipse reports 9 warnings related to classes defined in the JDK: 2 true positives (thus, it misses 43 real resource leaks) and 7 false positives. In "potential" leak mode, the analysis reports many more warnings. Thus, we triaged only the 180 warnings about JDK classes from the ZooKeeper benchmark. Among these were 3 true positives (it misses 9 real resource leaks) and 177 false positives . The most common cause of false positives was the unchangeable, default ownership transfer assumption at method invocations: the potential leak mode warned at each call that returns a resource-alias, such as `Socket#getInputStream`.

These results show that the Eclipse bug-finder is unsound in both modes: it misses 9 real warnings on ZooKeeper even in potential leak mode. The analysis is imprecise, too: in potential leak mode, the vast majority of warnings are false positives, and even in high-confidence mode its ratio of true to total warnings is higher than Plumber's (2/9 = 22% vs. 45/156 = 29%). Compared to our tool, the Eclipse analysis is much faster and easier to apply to legacy code (no annotations are required), but both less sound and less precise.

8.3.2 Grapple. Grapple [35] is a modern typestate-based resource leak analysis focused on high precision and (relative) scalability. Grapple models its alias and dataflow analyses as dynamic transitive-closure computations over graphs, and leverages novel path encodings and techniques from predecessor-system Graspan [30] to achieve both context- and path-sensitivity. Grapple contains four checkers, of which two are useful for detecting resource leaks. Unlike Plumber, Grapple is unsound; e.g., it performs a fixed bounded unrolling of loops to make path sensitivity tractable.

The Grapple authors have already evaluated their tool on earlier versions of the case study programs in section 8.1 [35]; Table 4 reproduces results from their paper. Based on these numbers, Grapple reports fewer false positives than Plumber. Unfortunately, we could not run Grapple's leak detection on our versions of the benchmarks, due to various hardcoded paths in its source code and limited documentation. Also, full details on Grapple's true and false-positive warnings are currently unavailable (we have requested them from the authors), so we cannot study the precision differences further. The run times in Table 4 show a stark difference with our work;

Plumber runs in minutes, whereas Grapple can take many hours. Further, to our best knowledge, Grapple is not modular, so any code modification will necessitate a full re-analysis. Plumber only needs to re-analyze modified code and possibly its dependents after a change, not unmodified dependencies.

Note that the TP and FP numbers are not perfectly comparable between a modular and a whole-program analysis. Our tool reports violations of a user-supplied specification (which takes effort to write but provides documentation benefits), so it can ensure that a library is correct for all possible clients. By contrast, Grapple checks a library in the context of one specific client.

9 LIMITATIONS AND THREATS TO VALIDITY

Like any tool that analyzes source code, Plumber only gives guarantees for code that it checks: the guarantee excludes native code, the implementation of unchecked libraries (such as the JDK), and code generated dynamically or by other annotation processors such as Lombok. Though the Checker Framework can handle reflection soundly [3], by default (and in our case studies) Plumber compromises this guarantee by assuming that objects returned by reflective invocations do not carry must-call obligations (but this behavior can be disabled). Within the bounds of a user-written warning suppression, Plumber assumes that 1) any errors issued can be ignored, and 2) all annotations written by the programmer are correct.

The results of our experiments may not generalize, compromising the external validity of the experimental results. Plumber may produce more false positives, require more annotations, or be more difficult to use if applied to other programs. It would be both easier and more useful to use Plumber from the inception of a project, rather than applying it after the code had already be written, as we did in our case studies: code could be annotated as it was written. This would ensure that the annotations match the intent of the programmers and would guide the programmers to a better design.

Like any practical system, it is possible that there might be bugs in the implementation of Plumber, or in the design of its analyses. We have mitigated this threat with code review and an extensive test suite for Plumber: 110 test classes containing 3,458 lines of non-comment, non-blank code. This test suite is publicly available and distributed with Plumber.

10 RELATED WORK

Most prior work on resource leak detection either uses program analysis to detect leaks or adds language features to prevent them. Here we focus on the most relevant work from these categories.

10.1 Analysis-based approaches

Static analysis. Tracker [29] performs inter-procedural dataflow analysis to detect resource leaks, with various additional features to make their tool practical, including issue prioritization and handling of wrapper types. Tracker avoids whole-program alias analysis to improve scalability, instead using a local, access-path-based approach. While Tracker scales well to large programs, it is deliberately unsound, unlike Plumber.

The Eclipse Compiler for Java includes a dataflow-based bug-finder for resource leaks [7]. Its analysis uses a fixed set of ownership heuristics and a fixed list of wrapper classes; unlike Plumber,

it is unsound. It is fast. Similar analyses—with similar trade-offs compared to Plumber—exist in other heuristic bug-finding tools, including SpotBugs [24], PMD [23], and Infer [15]. Section 8.3.1 experimentally evaluates the Eclipse analysis. Relda and Relda2 [14, 33] are unsound resource-leak detection approaches that are specialized to the Android framework with call graphs that model the framework's use of callbacks for releasing resources.

Typestate analysis [9, 25] can be used to find resource leaks. Grapple [35] is the most recent system to use this approach, leveraging a disk-based graph engine to achieve unprecedented scalability on a single machine. Compared to Plumber, Grapple is more precise but suffers from unsoundness and much longer run times. Section 8.3.2 gives a more detailed comparison to Grapple.

The CLOSER [6] automatically inserts Java code to dispose of resources when they are no longer "live" according to its dataflow analysis. Their approach requires an expensive alias analysis for soundness, as well as manually-provided aliasing specifications for linked libraries. Plumber uses accumulation analysis [8, 17] to achieve soundness without the need for a whole-program alias analysis. We build on the core analysis described in [17].

Dynamic analysis. Some approaches use dynamic analysis to ameliorate leaks. Resco [5] operates similarly to a garbage collector, tracking resources whose program elements have become unreachable. When a given resource (such as file descriptors) is close to exhaustion, the runtime runs Resco to clean up any resources of that type that are unreachable. With a static approach such as ours, a tool like Resco becomes unnecessary—leaks become impossible.

Automated test generation can also be used to try to detect resource leaks. For example, leaks in Android applications can be found by repeatedly running neutral—i.e. eventually returning to the same state—GUI actions [32, 34]. Other techniques focus on taking advantage of common misuse of the Android activity lifecycle [2]. Testing can only show the presence of bugs, not their absence; Plumber verifies that no resource leaks are present.

Data sets and surveys. The DroidLeaks benchmark [19] is a set of Android apps with known resource leak bugs. Unfortunately, it includes only the compiled apps. Plumber runs on source code, so we were unable to run Plumber on DroidLeaks. Ghanavati et al. [12] performed a detailed recent study of resource leaks and their repairs in Java projects, showing the pressing need for better tooling for resource leak prevention.

10.2 Language-based approaches

Ownership types and Rust. Ownership type systems [4] impose control over aliasing, which in turn enables guaranteeing other high-level properties, like absence of resource leaks. We do not discuss the vast literature on ownership type systems here (see Clarke et al. [4] for a survey). Instead, we focus on ownership types in Rust [18] as the most popular practical example of using ownership to prevent resource leaks.

For a detailed overview of ownership in Rust, see chapter 4 of [18]; we give a brief overview here. In Rust, ownership is used to manage both memory and other resources. Every value associated with a resource must have a *unique* owning pointer, and when an owning pointer's lifetime ends, the value is "dropped," ensuring all resources are freed. Rust's ownership type system statically prevents not only resource leaks, but also other important issues like "double-free" bugs (releasing a resource more than once) and "use-after-free" bugs (using a resource after it has been released). But, this power comes with a cost; to enforce uniqueness, non-owning pointers must be invalidated after an ownership transfer, and can no longer be used. Maintaining multiple usable pointers to a value requires use of language features like references and borrowing, and even then, borrowed pointers have restricted privileges.

Plumber has less power than Rust's ownership types; it cannot prevent double-free or use-after-free bugs. But, Plumber's lightweight ownership annotations impose *no* restrictions on aliasing; they simply aid the tool in identifying how a resource will be closed. Lightweight ownership is better suited to preventing resource leaks in existing, large Java code bases; adapting such programs to use a full Rust-style ownership type system would be impractical.

Other approaches. Java's try-with-resources construct [21] was discussed in section 1. Java also provides finalizer methods [13, Chapter 12], which execute before an object is garbage-collected, but they should not be used for resource management, as their execution may be delayed arbitrarily.

Compensation stacks [31] generalize C++ destructors and Java's try-with-resources, to avoid resource leak problems in Java. While compensation stacks make resource leaks less likely, they do not provide a guarantee that leaks will not occur, unlike Plumber.

An alternative to applying a typestate analysis to an existing program is to rewrite the program in a typestate-oriented programming language [1, 11]. The type systems of such languages are expressive enough to prevent resource leaks, but they impose restrictions on aliasing and a high type annotation burden, making translation of existing code impractical.

11 CONCLUSION

We have developed Plumber, a new, modular approach for detecting and preventing resource leaks in large-scale Java programs. Plumber is based on a sound core of an accumulation analysis augmented by three new features to handle aliasing patterns: lightweight ownership transfer, resource aliasing, and obligation creation by non-constructor methods. Our evaluation showed that Plumber is effective: it discovered 45 bugs in heavily-used, heavily-tested Java code that makes extensive use of resources, with analysis speed orders of magnitude faster than whole-program analysis, a false positive rate similar to a state-of-the-practice heuristic bug-finder, and a manageable annotation burden of 1 annotation per 2,000 lines of code. In future work, we plan to develop improved inference techniques for lightweight ownership annotations. Since these annotations can be added anywhere without impacting soundness, genetic search (using the warnings emitted by Plumber as the fitness function) and machine-learning techniques could be used to find an optimal placement of ownership annotations.

REFERENCES

[1] Jonathan Aldrich, Joshua Sunshine, Darpan Saini, and Zachary Sparks. 2009. Typestate-oriented programming. In *OOPSLA Companion: Object-Oriented Programming Systems, Languages, and Applications*. Orlando, FL, USA, 1015–1022.

[2] Domenico Amalfitano, Vincenzo Riccio, Porfirio Tramontana, and Anna Rita Fasolino. 2020. Do memories haunt you? An automated black box testing approach

for detecting memory leaks in android apps. *IEEE Access* 8 (2020), 12217–12231.

[3] Paulo Barros, René Just, Suzanne Millstein, Paul Vines, Werner Dietl, Marcelo d'Amorim, and Michael D. Ernst. 2015. Static analysis of implicit control flow: Resolving Java reflection and Android intents. In *ASE 2015: Proceedings of the 30th Annual International Conference on Automated Software Engineering*. Lincoln, NE, USA, 669–679.

[4] Dave Clarke, Johan Östlund, Ilya Sergey, and Tobias Wrigstad. 2013. Ownership Types: A Survey. In *Aliasing in Object-Oriented Programming. Types, Analysis and Verification*. Springer, Berlin, Heidelberg.

[5] Ziying Dai, Xiaoguang Mao, Yan Lei, Xiaomin Wan, and Kerong Ben. 2013. Resco: Automatic collection of leaked resources. *IEICE TRANSACTIONS on Information and Systems* 96, 1 (2013), 28–39.

[6] Isil Dillig, Thomas Dillig, Eran Yahav, and Satish Chandra. 2008. The CLOSER: automating resource management in Java. In *International symposium on Memory management*. 1–10.

[7] Eclipse developers. 2020. Avoiding resource leaks. https://help.eclipse.org/2020-12/index.jsp?topic=%2Forg.eclipse.jdt.doc.user%2Ftasks%2Ftask-avoiding_resource_leaks.htm&cp%3D1_3_9_3. Accessed 3 February 2021.

[8] Manuel Fähndrich and K. Rustan M. Leino. 2003. Heap Monotonic Typestates. In *IWACO 2003: International Workshop on Aliasing, Confinement and Ownership in object-oriented programming*. Darmstadt, Germany.

[9] Stephen J. Fink, Eran Yahav, Nurit Dor, G. Ramalingam, and Emmanuel Geay. 2008. Effective typestate verification in the presence of aliasing. *ACM TOSEM* 17, 2, Article Article 9 (2008), 34 pages.

[10] Jeffrey S. Foster, Manuel Fähndrich, and Alexander Aiken. 1999. A theory of type qualifiers. In *PLDI '99: Proceedings of the ACM SIGPLAN '99 Conference on Programming Language Design and Implementation*. Atlanta, GA, USA, 192–203.

[11] Ronald Garcia, Éric Tanter, Roger Wolff, and Jonathan Aldrich. 2014. Foundations of Typestate-Oriented Programming. *ACM Trans. Program. Lang. Syst.* 36, 4 (2014).

[12] Mohammadreza Ghanavati, Diego Costa, Janos Seboek, David Lo, and Artur Andrzejak. 2020. Memory and resource leak defects and their repairs in Java projects. *Empirical Software Engineering* 25, 1 (2020), 678–718.

[13] James Gosling, Bill Joy, Guy L. Steele, Gilad Bracha, and Alex Buckley. 2014. *The Java Language Specification, Java SE 8 Edition*. Addison-Wesley Professional.

[14] Chaorong Guo, Jian Zhang, Jun Yan, Zhiqiang Zhang, and Yanli Zhang. 2013. Characterizing and detecting resource leaks in Android applications. In *Automated Software Engineering (ASE)*. IEEE, 389–398.

[15] Infer developers. 2021. Resource leak in Java. https://fbinfer.com/docs/checkers-bug-types#resource-leak-in-java. Accessed 4 February 2021.

[16] JetBrains. 2020. List of Java Inspections. https://www.jetbrains.com/help/idea/list-of-java-inspections.html#resource-management. Accessed 5 February 2021.

[17] Martin Kellogg, Manli Ran, Manu Sridharan, Martin Schäf, and Michael D. Ernst. 2020. Verifying Object Construction. In *ICSE 2020, Proceedings of the 42nd International Conference on Software Engineering*. Seoul, Korea.

[18] Steve Klabnik and Carol Nichols. 2018. *The Rust Programming Language*. https://doc.rust-lang.org/1.50.0/book/

[19] Yepang Liu, Jue Wang, Lili Wei, Chang Xu, Shing-Chi Cheung, Tianyong Wu, Jun Yan, and Jian Zhang. 2019. DroidLeaks: a comprehensive database of resource leaks in Android apps. *Empirical Software Engineering* 24, 6 (2019), 3435–3483.

[20] David Lo, Nachiappan Nagappan, and Thomas Zimmermann. 2015. How practitioners perceive the relevance of software engineering research. In *ESEC/FSE 2015: The 10th joint meeting of the European Software Engineering Conference (ESEC) and the ACM SIGSOFT Symposium on the Foundations of Software Engineering (FSE)*. Bergamo, Italy.

[21] Oracle. 2020. The try-with-resources Statement (The Java Tutorials). https://docs.oracle.com/javase/tutorial/essential/exceptions/tryResourceClose.html. Accessed 24 February 2021.

[22] Matthew M. Papi, Mahmood Ali, Telmo Luis Correa Jr., Jeff H. Perkins, and Michael D. Ernst. 2008. Practical pluggable types for Java. In *ISSTA 2008, Proceedings of the 2008 International Symposium on Software Testing and Analysis*. Seattle, WA, USA, 201–212.

[23] PMD developers. 2021. CloseResource. https://pmd.github.io/pmd-6.31.0/pmd_rules_java_errorprone.html#closeresource. Accessed 4 February 2021.

[24] SpotBugs developers. 2021. OBL: Method may fail to clean up stream or resource. https://spotbugs.readthedocs.io/en/latest/bugDescriptions.html#obl-method-may-fail-to-clean-up-stream-or-resource-obl-unsatisfied-obligation. Accessed 4 February 2021.

[25] Robert E. Strom and Shaula Yemini. 1986. Typestate: A programming language concept for enhancing software reliability. *IEEE TSE* SE-12, 1 (Jan. 1986), 157–171.

[26] Bjarne Stroustrup. 1994. 16.5, Resource Management. In *The design and evolution of C++*. Pearson Education India, 388–389.

[27] The Apache Hadoop developers. 2018. StorageInfo.java. https://github.com/apache/hadoop/blob/aa96f1871bfd858f9bac59cf2a81ec470da649af/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/StorageInfo.java#L246. Accessed 22 February 2021.

[28] The Apache ZooKeeper developers. 2020. Learner.java. https://github.com/apache/zookeeper/blob/c42c8c94085ed1d94a22158fbdfe2945118a82bc/zookeeper-server/src/main/java/org/apache/zookeeper/server/quorum/Learner.java#L465. Accessed 24 February 2021.

[29] Emina Torlak and Satish Chandra. 2010. Effective interprocedural resource leak detection. In *International Conference on Software Engineering (ICSE)*. 535–544.

[30] Kai Wang, Aftab Hussain, Zhiqiang Zuo, Guoqing (Harry) Xu, and Ardalan Amiri Sani. 2017. Graspan: A Single-machine Disk-based Graph System for Interprocedural Static Analyses of Large-scale Systems Code. In *Architectural Support for Programming Languages and Operating Systems (ASPLOS)*.

[31] Westley Weimer and George C Necula. 2004. Finding and preventing run-time error handling mistakes. In *Object-oriented programming, systems, languages, and applications (OOPSLA)*. 419–431.

[32] Haowei Wu, Yan Wang, and Atanas Rountev. 2018. Sentinel: generating GUI tests for Android sensor leaks. In *International Workshop on Automation of Software Test (AST)*. IEEE, 27–33.

[33] Tianyong Wu, Jierui Liu, Xi Deng, Jun Yan, and Jian Zhang. 2016. Relda2: an effective static analysis tool for resource leak detection in Android apps. In *Automated Software Engineering (ASE)*. IEEE, 762–767.

[34] Hailong Zhang, Haowei Wu, and Atanas Rountev. 2016. Automated test generation for detection of leaks in Android applications. In *International Workshop on Automation of Software Test (AST)*. 64–70.

[35] Zhiqiang Zuo, John Thorpe, Yifei Wang, Qiuhong Pan, Shenming Lu, Kai Wang, Guoqing Harry Xu, Linzhang Wang, and Xuandong Li. 2019. Grapple: A graph system for static finite-state property checking of large-scale systems code. In *EuroSys*. 1–17.